



Apache Spark

SAMEER MAHAJAN



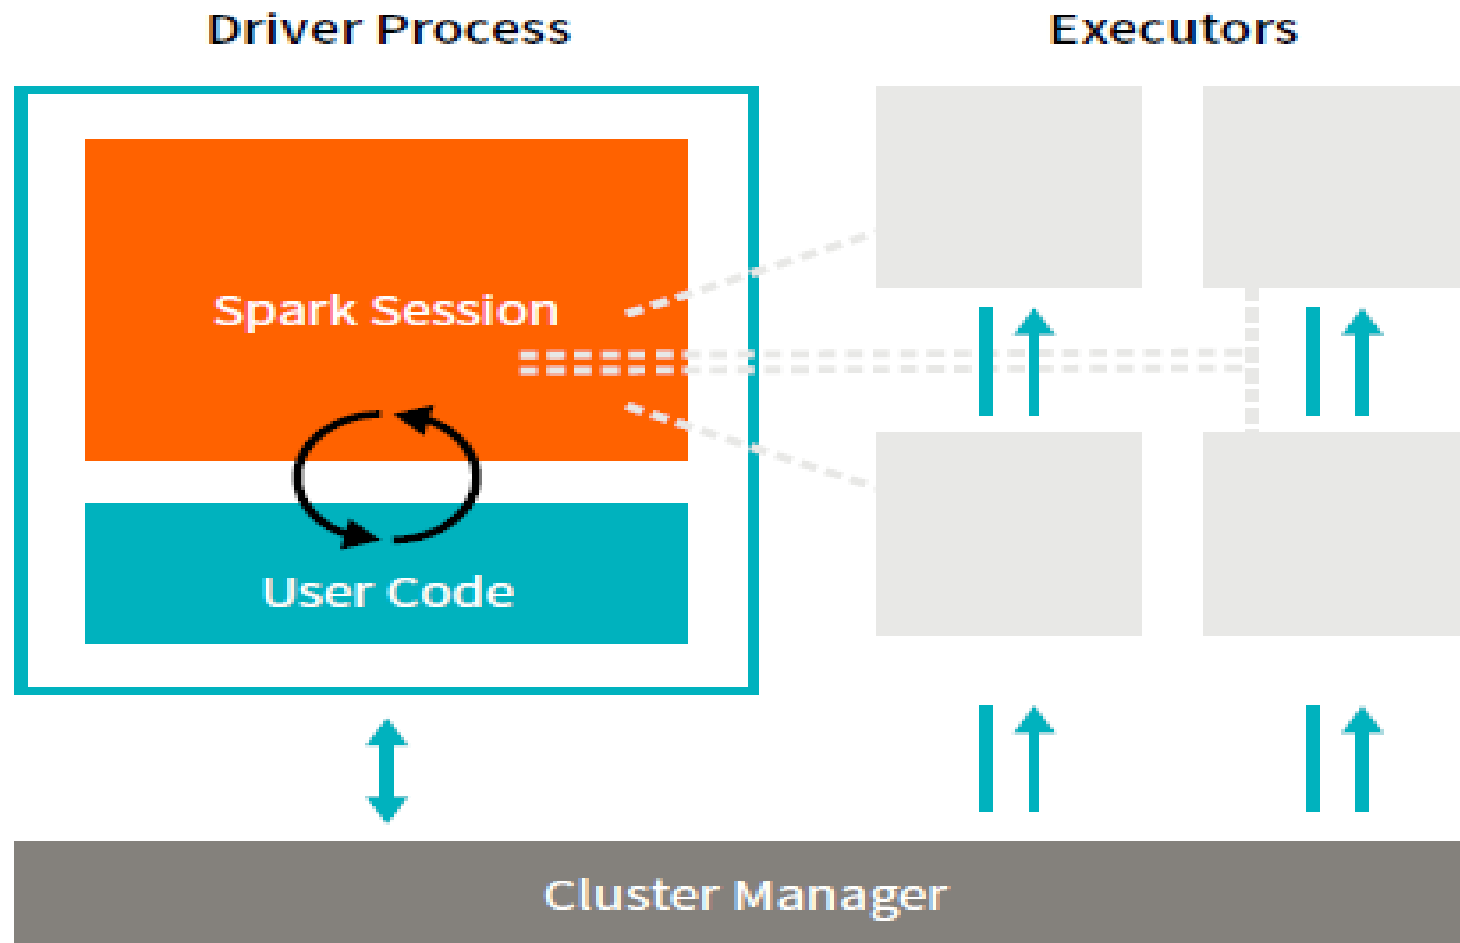
History

- 2009 research project in UC Berkeley
- Learnt from Hadoop
 - Cluster computing
 - Multi pass over data
- API based on functional programming that could succinctly express multi-step application
- in-memory data sharing across computation steps
- Version 1: Batch Processing only, Version 2: Interactive
- Shark: SQL queries, 2011
- AMPlab contributed Spark to Apache and started Databricks
- Apache Spark 1.0 in 2014 (Spark SQL API) and 2.0 in 2016
- Structured Streaming, 2017

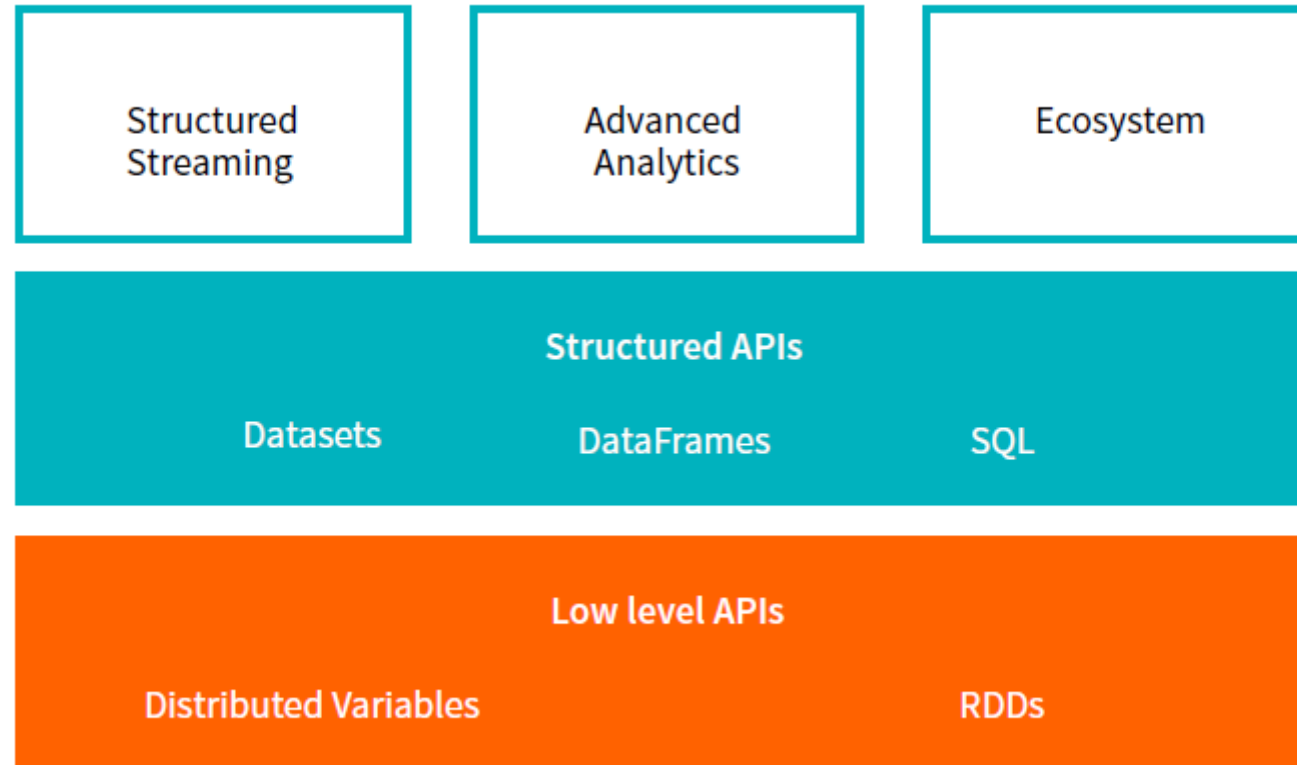
Key Benefits

- Parallelism on a cluster
- Most actively developed open source engine
- Wide support for programming languages (Python, Java, Scala, R)
- Unified analytics compute engine for loading, querying, streaming, ML
- Consistent, composable APIs
- Runs from a laptop to a cluster of thousands of servers
- Works with s3, Hadoop, Cassandra, Kafka etc.
- Large number of internal and external (<http://spark-packages.org>) libraries

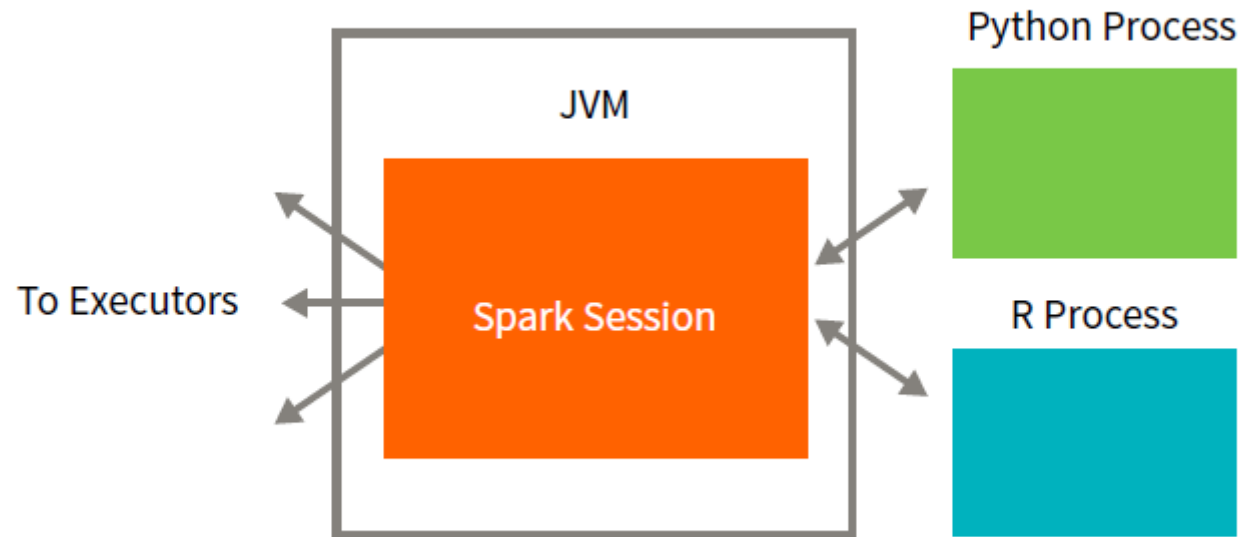
High level architecture



High level view



Language support



Data Structures and Some Basic Concepts

- Data Structures
 - Dataframe
 - Dataset
 - SQL Table
 - Resilient Distributed Dataset (RDD)
- Some Basic Concepts
 - Partitions
 - Transformations
 - Lazy evaluation
 - Action: triggers the computation

Getting started

- Operate at spark dataframe level as much as possible
- Spark application is controlled through a driver process
- Driver process manifests as spark session
- One to one mapping between instance of spark session and spark application

Basics

- `spark.read` to read csv
- `spark.createDataFrame`
 - define schema, specify values
- `.repartition(1).write.csv('location')` : write spark DF to a single file

Some conversions

- `.rdd` to convert to rdd (`.rdd.map (lambda row: row[o])` to convert a column to rdd)
- `rdd.toDF()` : to convert rdd to spark DF
- `toPandas()` : convert spark DF to pandas DF
- `SQLContext(sc). createDataFrame(pandas_df)` : convert pandas DF to spark DF
- `df.createOrReplaceTempView('sql_table')` : convert spark DF into sql table

Apache Spark Basics

define schema

```
In [ ]: from pyspark.sql.types import StructType
from pyspark.sql.types import StructField
from pyspark.sql.types import StringType
from pyspark.sql.types import FloatType

df_schema = StructType([StructField("Wave_Name", StringType(), True), StructField("Manufacturer", StringType(), True),
                           StructField("Gender", StringType(), True), StructField("Task", StringType(), True),
                           StructField("Label", FloatType(), True), StructField("Prediction", FloatType(), True)
                           ])
```

create data frame using specified schema and empty data

```
In [ ]: df = spark.createDataFrame(sc.emptyRDD(), df_schema)
```

reading csv file into data frame

```
In [ ]: df = spark\
    .read\
    .option("inferSchema", "true")\
    .option("header", "true")\
    .csv('<file path>')
```

Machine learning libraries

- ml
 - newer implementation
 - still under development with some experimental classes
 - works on spark Dataframe
 - watch out for 'Experimental'
- mllib
 - older implementation
 - works on RDD

handle missing values and extract day of week

```
In [ ]: prepped_df = df\  
        .na.fill(0)\  
        .withColumn("day_of_week", date_format(col("InvoiceDate"), "EEEE"))\  
        .coalesce(5)
```

split data into train and test data sets

```
In [ ]: train_df, test_df = prepped_df.randomSplit([0.8, 0.2], seed=12345)
```

convert string column values into numeric indexers

```
In [ ]: from pyspark.ml.feature import StringIndexer  
  
indexer = StringIndexer()\br/>        .setInputCol("day_of_week")\  
        .setOutputCol("day_of_week_index")
```

use one hot encoding

```
In [ ]: from pyspark.ml.feature import OneHotEncoder  
  
encoder = OneHotEncoder()\br/>        .setInputCol("day_of_week_index")\  
        .setOutputCol("day_of_week_encoded")
```

Parallel execution

- spark manages splitting into and aggregating results from individual tasks distributed on workers
- `sc.parallelize` and `map lambda`
 - parallel task runs on a worker (not driver) node
 - limitations on what can be passed to and used inside the parallel task
- python threads

using parallelize

```
In [ ]: tasks = []
        for manufacturer in manufacturers:
            for gender in genders:
                for speech_task in speech_tasks:
                    for lower_threshold in lower_thresholds:
                        for higher_threshold in higher_thresholds:
                            tasks = tasks + [(manufacturer, gender, speech_task, lower_threshold, higher_threshold)]

        tasksRDD = sc.parallelize(tasks, numSlices = len(tasks))

        cfc = tasksRDD.map(lambda alpha: compute_ML_grid_for_a_cohort(alpha[0], alpha[1], alpha[2], alpha[3], alpha[4]))
        cfc.cache()
        cfc.count()
```

using python threads for parallelization

```
In [ ]: from threading import Thread
        num_cohorts = len(manufacturers) * len(genders) * len(speech_tasks) * len(lower_thresholds) * len(higher_thresholds)
        tasks = [None] * num_cohorts
        results = [None] * num_cohorts

        i = 0
        for manufacturer in manufacturers:
```

References

- <http://spark.apache.org/>
- <https://pages.databricks.com/gentle-intro-spark.html>
- <http://shop.oreilly.com/product/0636920034957.do>
- Spark Python API Documentation:
<https://spark.apache.org/docs/latest/api/python/index.html>